




# Using Kafka Streams to Build a Data Pipeline for the Hospitality Industry

September 17, 2024



## COLLECT

Multi source data  
connectivity & ingestion

## UNIFY / MERGE

Data processing engine

## STORE




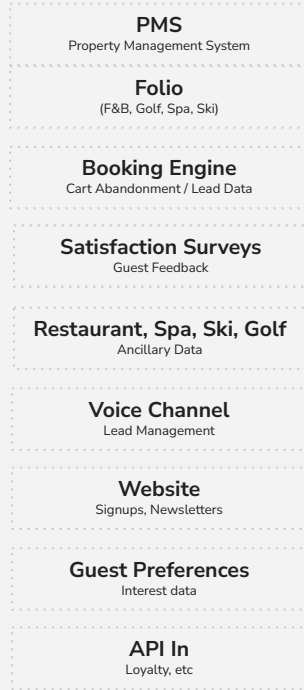
Store and search data

## TARGET

Define rules &  
create segments

## ACTIVATE

Connect & send data  
to external systems



EMAIL

VOICE

MESSAGING

SMS

WEB

API

**Systems & Data Integrations**  
Turnkey data activation and data flow  
management running on secure AWS Cloud

**Advanced Profile Synthesis**  
Merge, dedupe, unify, organize  
and secure clean guest data

**#1 Hospitality CRM**  
Storage and visualization of all  
hotel and brand Rich Guest Profiles

**Advanced Data Segmentation** Set  
up segments and audiences for  
automated personalization

**Omnichannel Activation**  
Powering personalized guest  
lifecycle communications on all  
your direct channels natively

**450+**

REVINATORS  
WORLDWIDE  
(US, EMEA, APAC)

**\$150M+**

INVESTED IN  
HOTELIER SUCCESS,  
BACKED BY TOP  
GLOBAL INVESTORS



**#1 CDP/CRM**

THE ULTIMATE  
OMNI-CHANNEL DIRECT  
BOOKING PLATFORM  
FOR HOTELS

**800M+**

RICH GUEST  
PROFILES UNDER  
MANAGEMENT



**REVINATE**



**60 NPS**

HOTELIERS LOVE  
OUR PURPOSE  
BUILT PLATFORM &  
EXPERT SERVICE

**\$12.8B+**

LIFETIME DIRECT  
REVENUES GENERATED  
FOR REVINATE  
CUSTOMERS



## 6 Years Ago

- Monolith
- Java with Spring Boot
- Self-managed Kafka
- Spring Kafka
- Limited Kafka knowledge
- RabbitMQ



## Today

- Microservices
- Vanilla Java
- Hosted Kafka
- Kafka Streams
- More Kafka knowledge

## Data


- Accounts (hotels)
- Reservations (room, spa, golf, restaurant)
- Guests (profiles, preferences, feedback)

## Processing

- Multiple integration methods (webhooks, APIs, FTP, email)
- Realtime and batch workloads (duplicate source system messages are common)
- At-least-once
- Atomic messages, but with complex schema
- Order isn't always reliable from source systems



# 1. Supporting Diverse Message Types



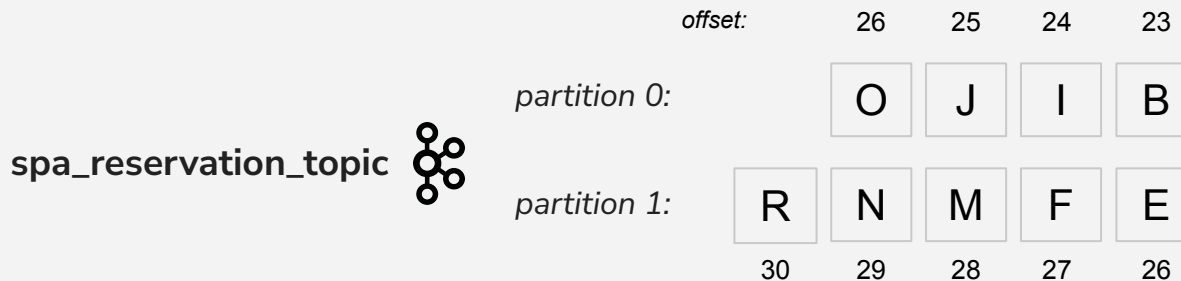
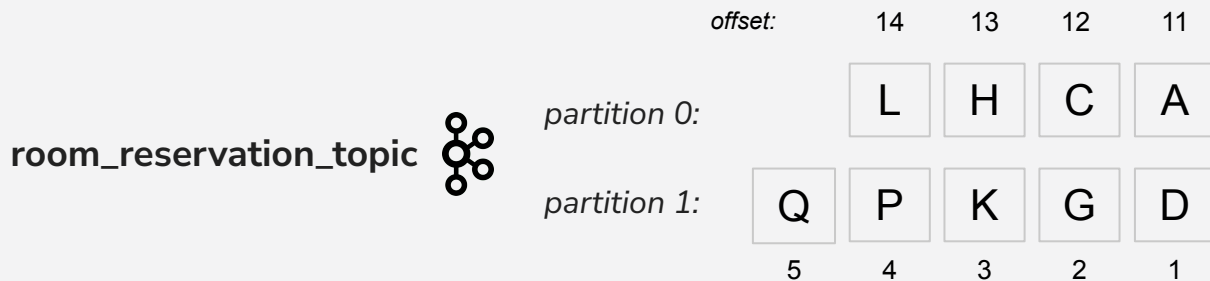
## ONE POTENTIAL APPROACH

*"The common wisdom seems to be: put all events of the same type in the same topic, and use different topics for different event types" <sup>1</sup>*



## ORDERING CONSIDERATIONS

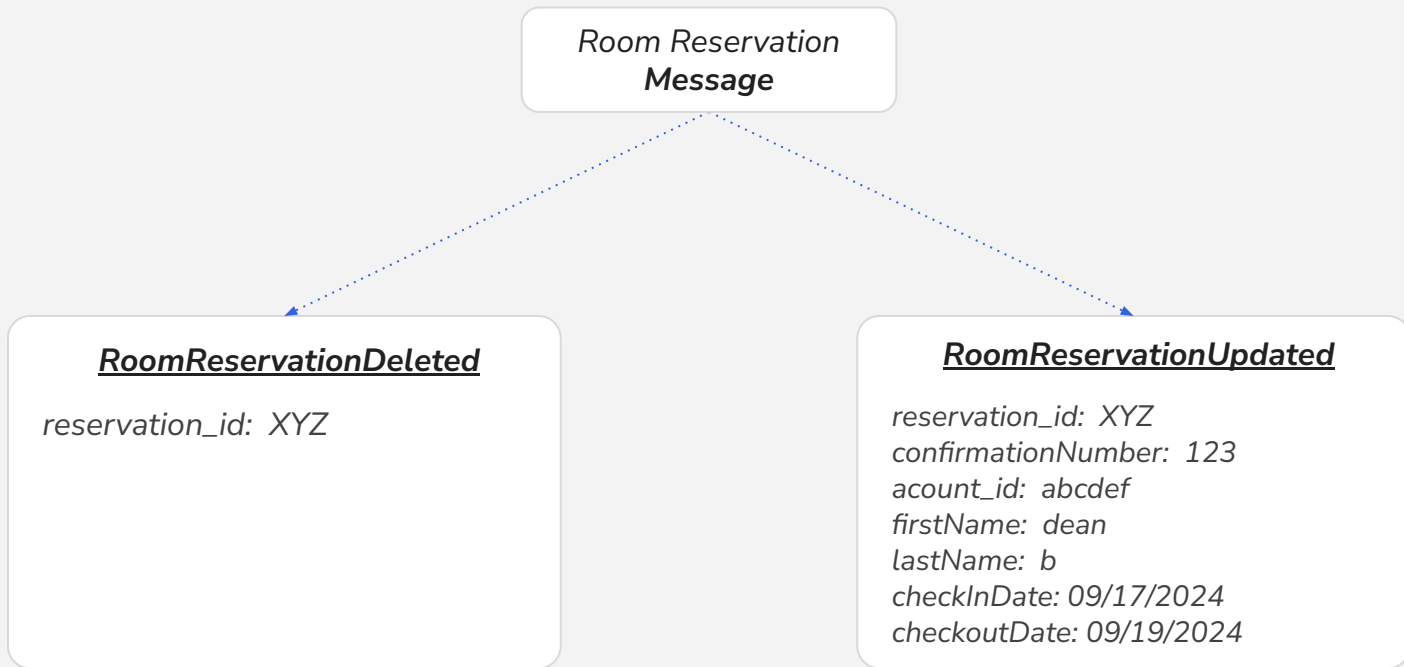
*"Kafka only provides a total order over records within a partition, not between different partitions in a topic" <sup>1</sup>*





## Problem 1

Messages with different types can't be ordered because they are on different topics




## Problem 2

Multiple message types can be a challenge to maintain when they share common components



### Problem 3

- Services may need to support multiple serdes for input and output topics
- Adding new schemas requires the addition of new topics

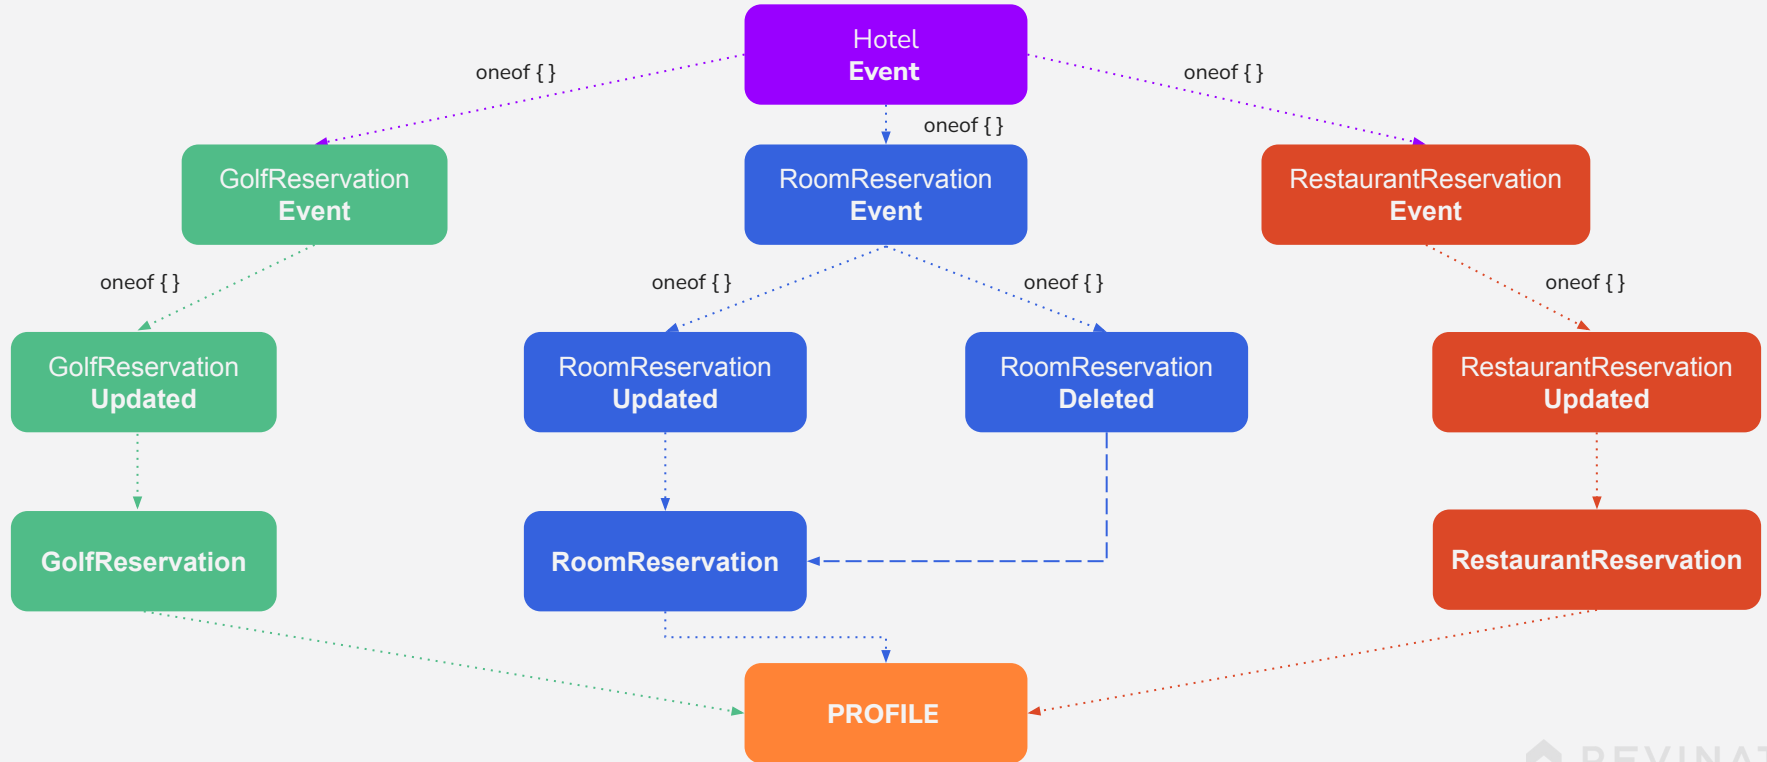


## Goal

- New message types can be supported without creating new topics
- Message order can be guaranteed *across* message types
- Consistent processing of shared components
- One serde for all

## Solution

Create a generalized schema that can be used for all topics




## Disadvantage

```
void process(HotelEvent hotelEvent) {  
    switch(hotelEvent.eventType()) {  
        case roomReservation -> {  
            switch (roomRervation.getType()) {  
                case roomReservationUpdate -> {...}  
                case roomReservationDelete -> {...}  
            }  
        }  
        case profile -> {...}  
        case golfReservation -> {...}  
        case spaReservation -> {...}  
    }  
}
```



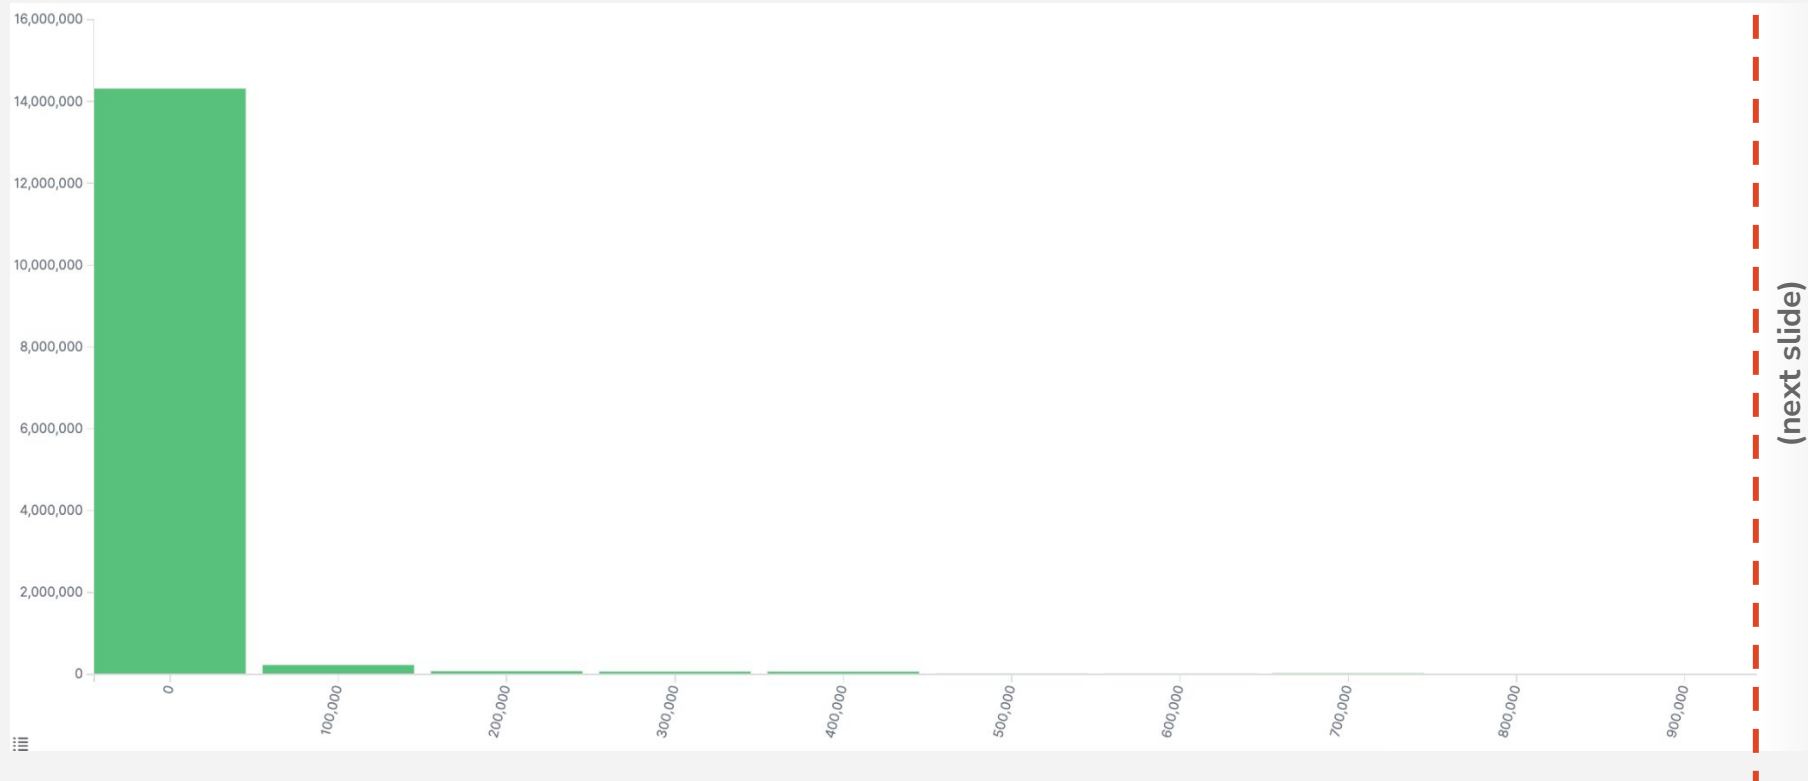
## 2. Handling Large Messages



NEARLY ALL OF OUR MESSAGES ARE UNDER 1MB

98.8% of messages are under 1MB (default Kafka limit)

1MB  
under | over



Source: Revinate logging infrastructure (sample of 15 million messages)




SOME MESSAGES ARE OVER 1MB

1MB

Only 1.2% of messages are between 1MB and 10MB...but we must support them

10MB  
under over




Source: Revinate logging infrastructure (sample of 15 million messages)

SOME MESSAGES ARE WAY OVER 1MB

10MB

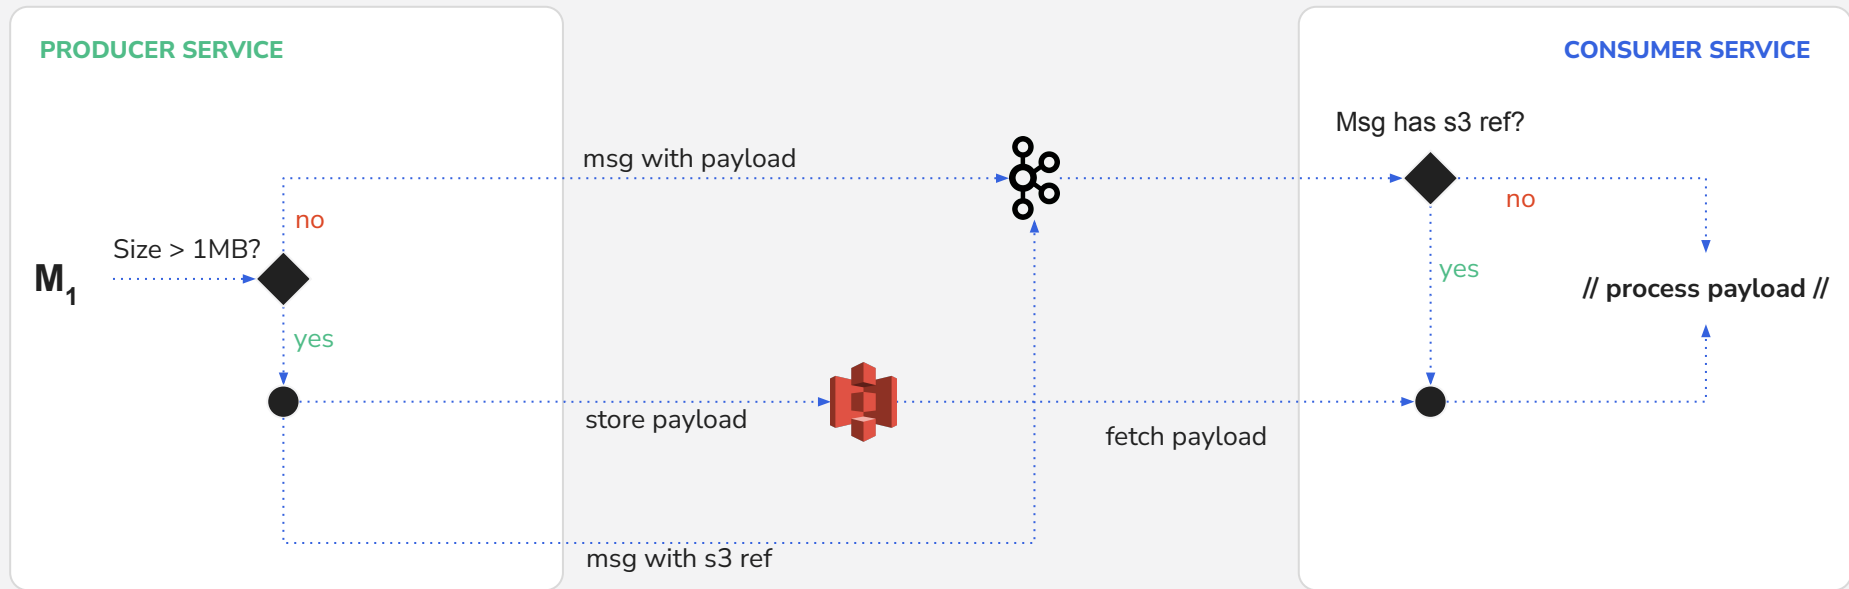
Barely any messages (0.005%) are over 10MB...but we must support them, too



**Source:** Revinate logging infrastructure (sample of 15 million messages)

## Solution

Use the Claim-Check pattern <sup>1</sup>




## Some things to consider

- Compression will affect message size (`message.max.bytes`)

"As with all byte sizes specified on the broker, this configuration deals with compressed message size, which means that producers can send messages that are much larger than this value uncompressed, provided they compress down to under the configured `message.max.bytes` size."<sup>1</sup>

- Message schema changes
- S3 message payload retention
- Alternative approaches
  - Custom consumer/producer
  - *Likafka*<sup>2</sup>

# 3. Deduplicating Message Payloads



## Problem

- Message stream contains messages that are ***duplicates*** of messages that were previously processed
  - May result in erroneous duplicate processing errors
  - Unnecessary processing (waste of resources)

## Cause

- 3rd party processing beyond our control
- Intentional data replay

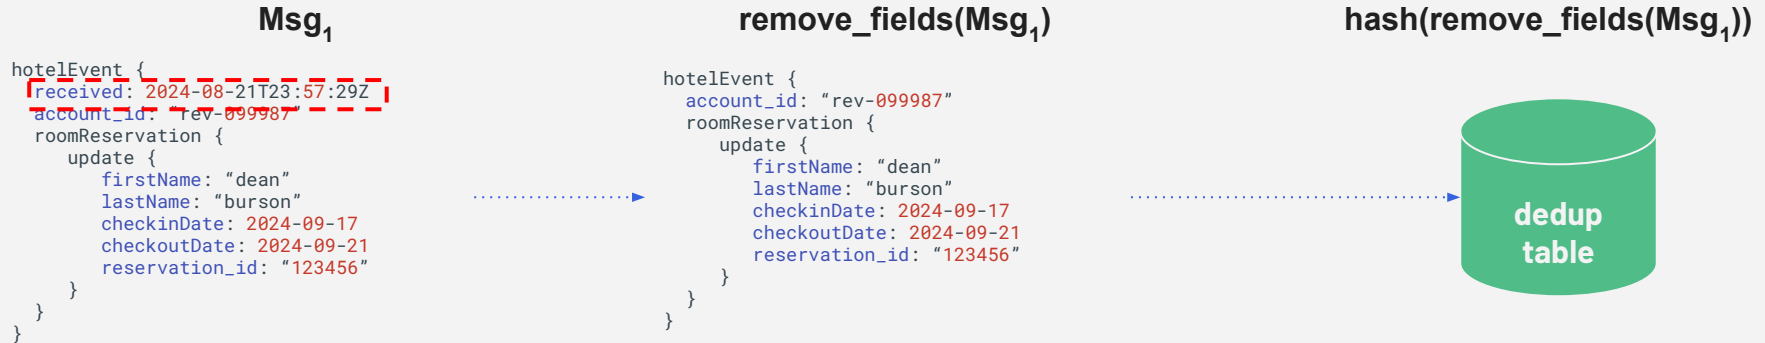
## Goal

- Identify duplicate messages
  - Redirect duplicates to a duplicate topic
  - Ignore duplicates

A message is a duplicate when (significant) fields are equivalent to fields of a previously processed message.

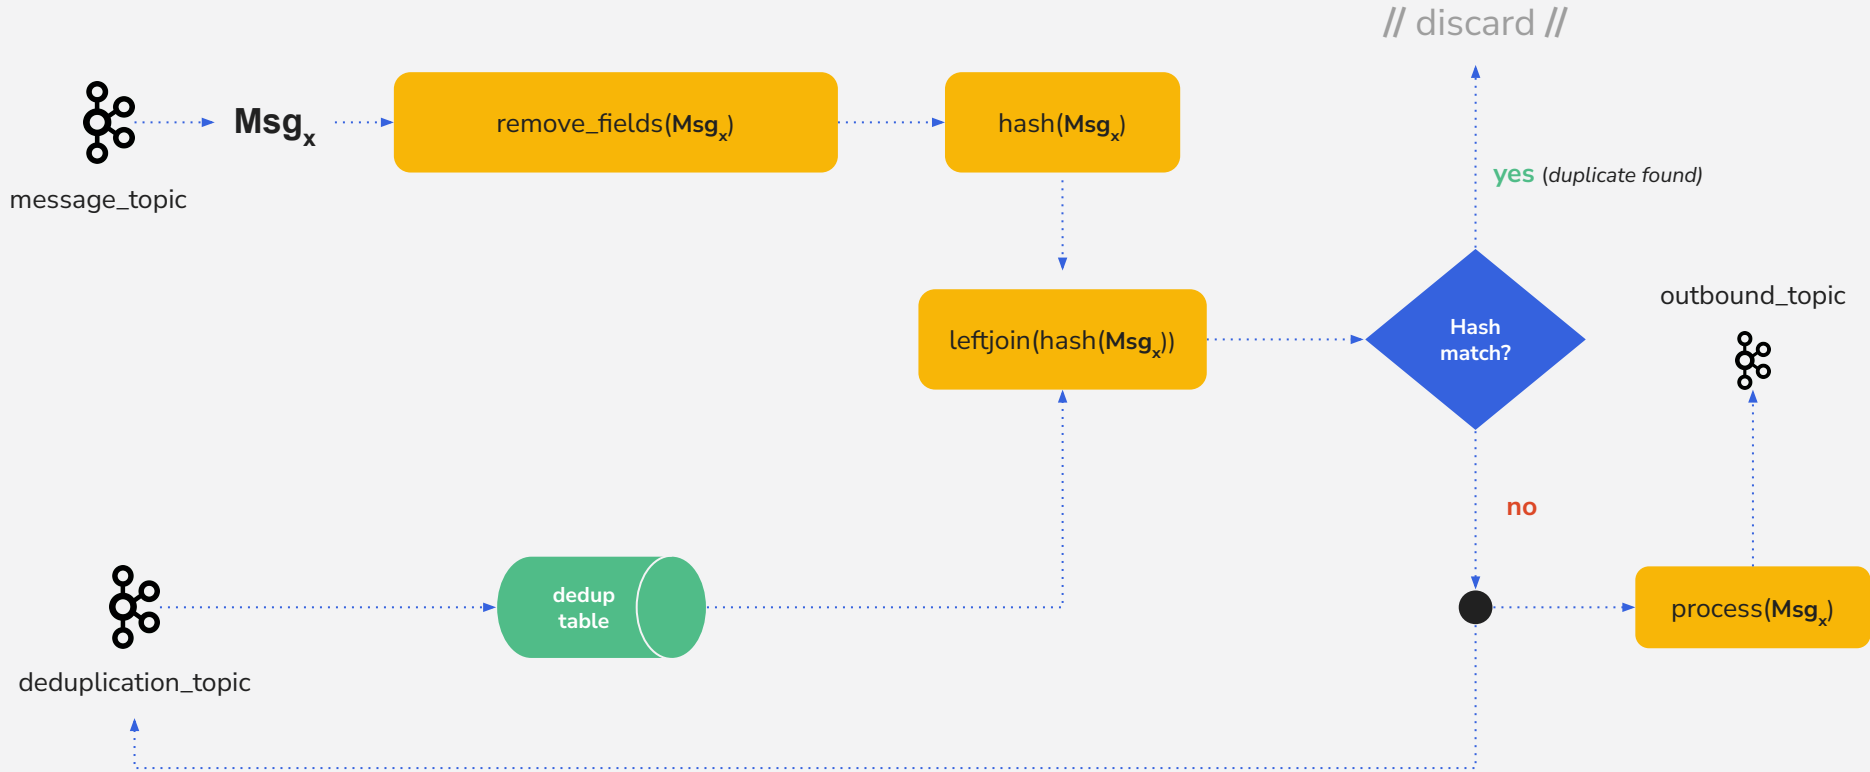
<b>Msg<sub>1</sub></b>	<b>Msg<sub>2</sub></b>
<pre>hotelEvent {   received: 2024-08-21T23:57:29Z   account_id: "rev-099987"   roomReservation {     update {       name: "dean"       checkinDate: 2024-09-17       checkoutDate: 2024-09-21       reservation_id: "123456"     }   } }</pre>	<pre>hotelEvent {   received: 2024-08-22T21:43:12Z   account_id: "rev-099987"   roomReservation {     update {       name: "dean"       checkinDate: 2024-09-17       checkoutDate: 2024-09-21       reservation_id: "123456"     }   } }</pre>

Remove non-significant fields, hash the result, and add to cache table





## Solution




## Some things to consider

- Dedup table structure
  - Message hash (Long as Key)
  - Current timestamp (for table maintenance)
- How long are messages retained in dedup table?
  - Hash cache may grow forever
  - How is time-to-live function implemented on dedup table
- Hash algorithm
- Will duplicate messages have same key?
  - If so, use partition KTable for dedup
  - If not, use GlobalKTable for dedup




# 4. Prioritizing Messages




## Problem

Topic contains messages whose processing priority is not reflected in partition order or by timestamp

**High-priority** message (E) needs to be processed ahead of messages received before



**Low-priority** messages (V, W, X, Y) delay processing of messages received after



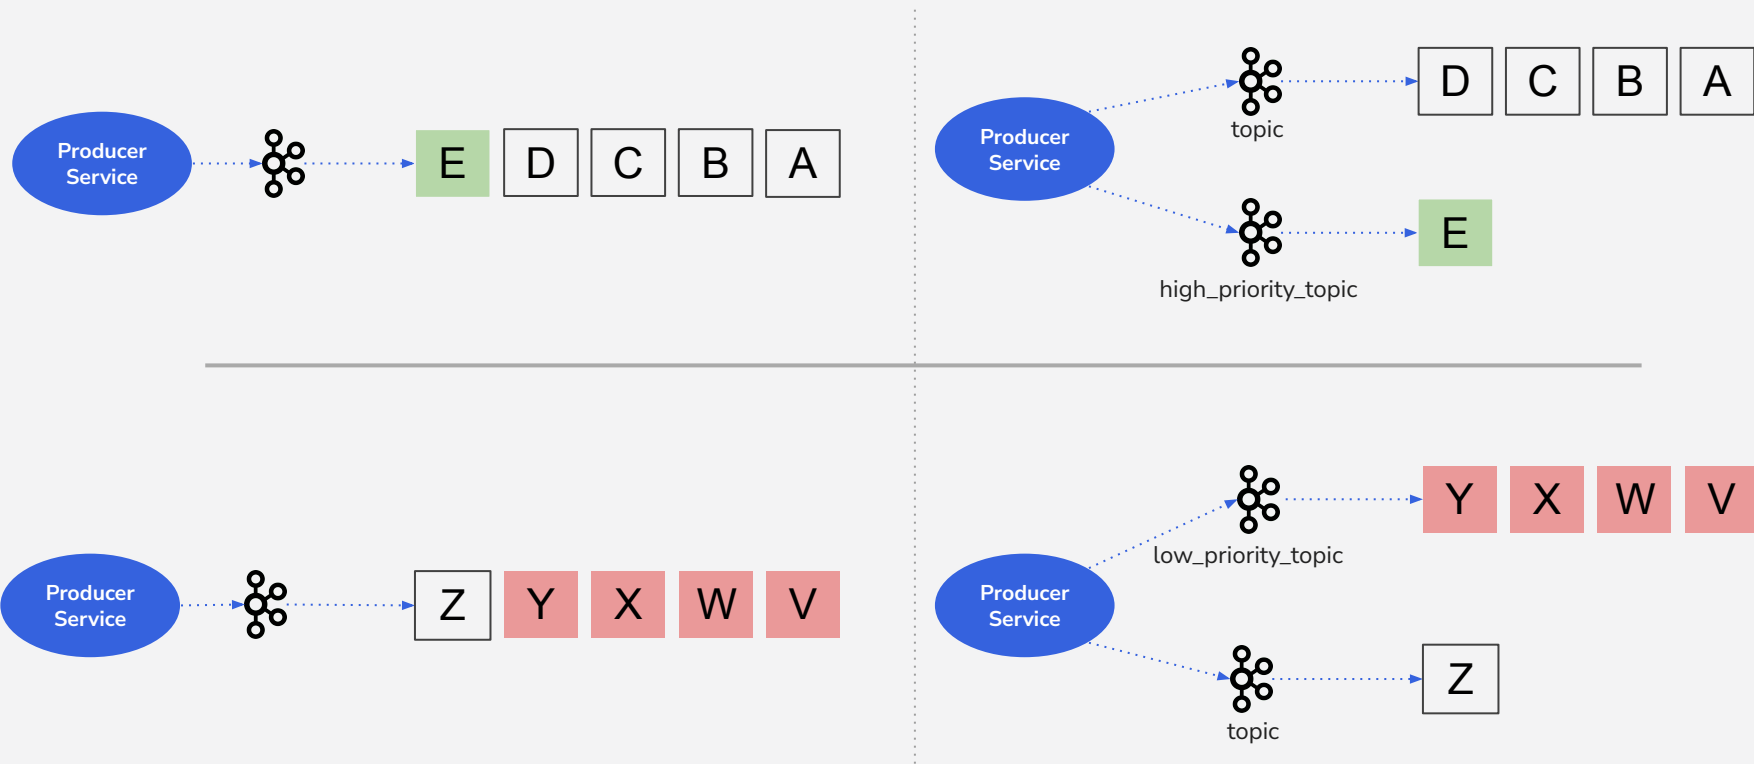
## Goal

- Identify high-priority and low-priority messages
- Ensure that high-priority messages are processed as soon as possible
- Ensure that low-priority message do not block or impede processing of higher priority message

## Steps

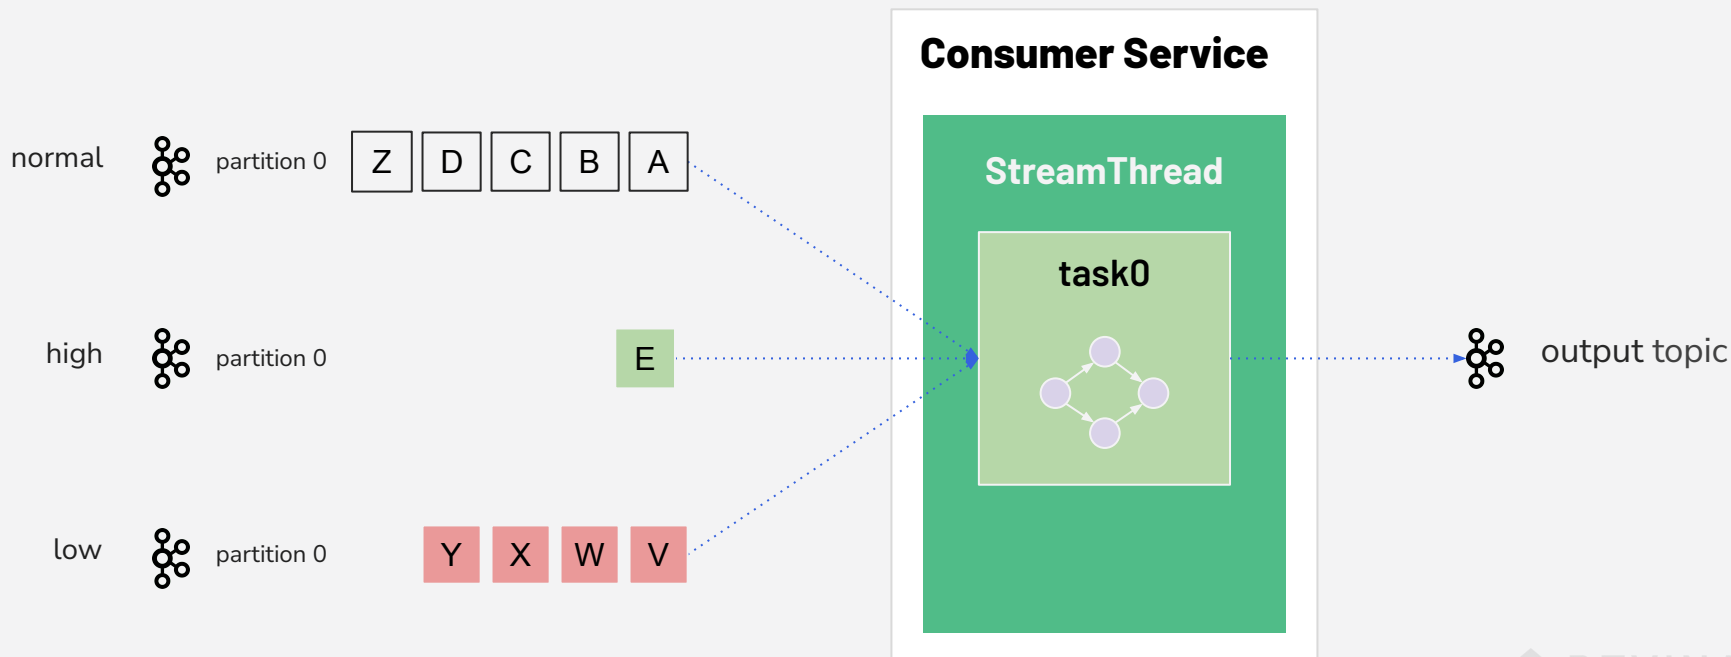
- Sort high/low priority messages into high/low priority topics
- Consume high/low priority topics using sub-topologies

ONE IDEA: USE SEPARATE TOPICS

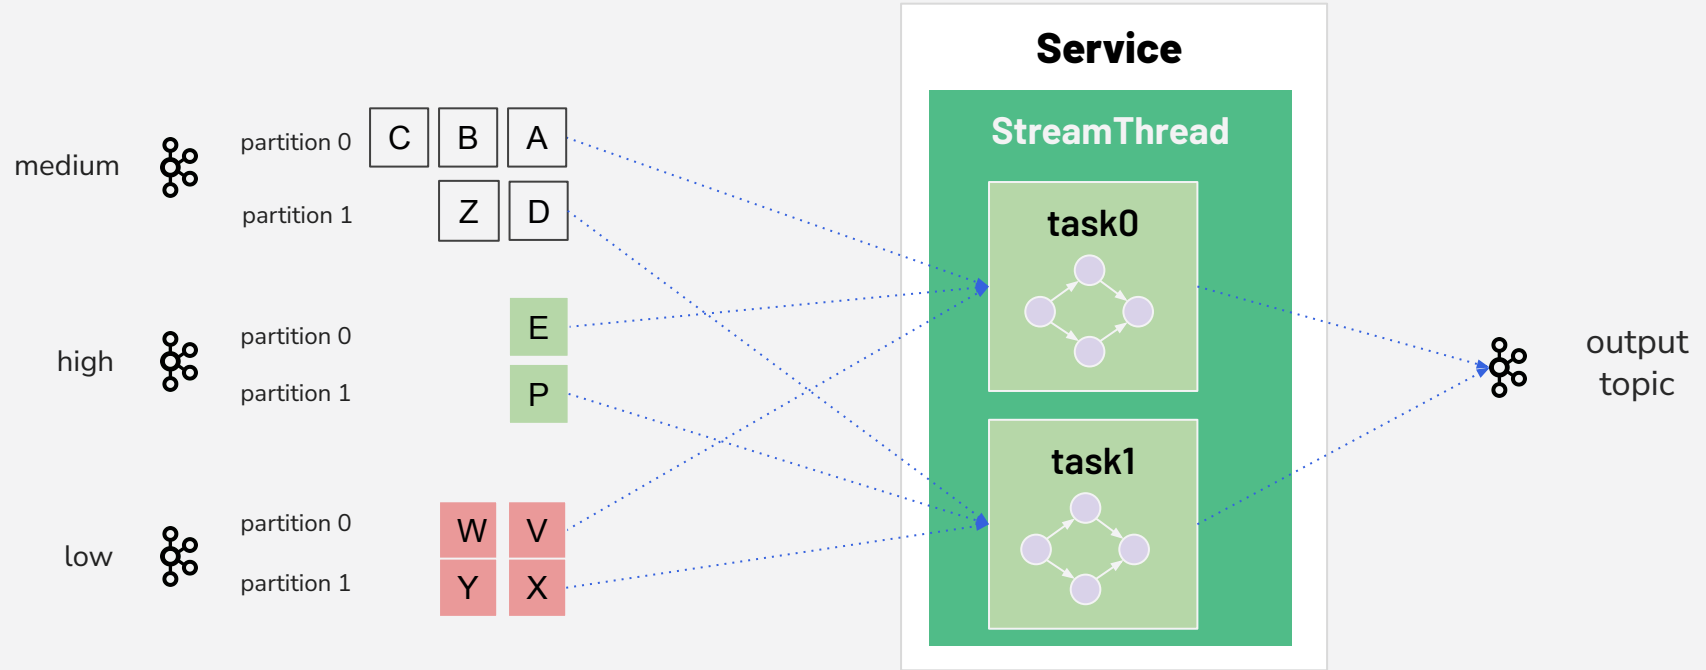


Sorting high/low priority messages ensures that messages are processed with minimal blocking (round-robin partition consumption), BUT...

*"In Kafka Streams, the earliest timestamp across all partitions is chosen first for processing" <sup>1</sup>*



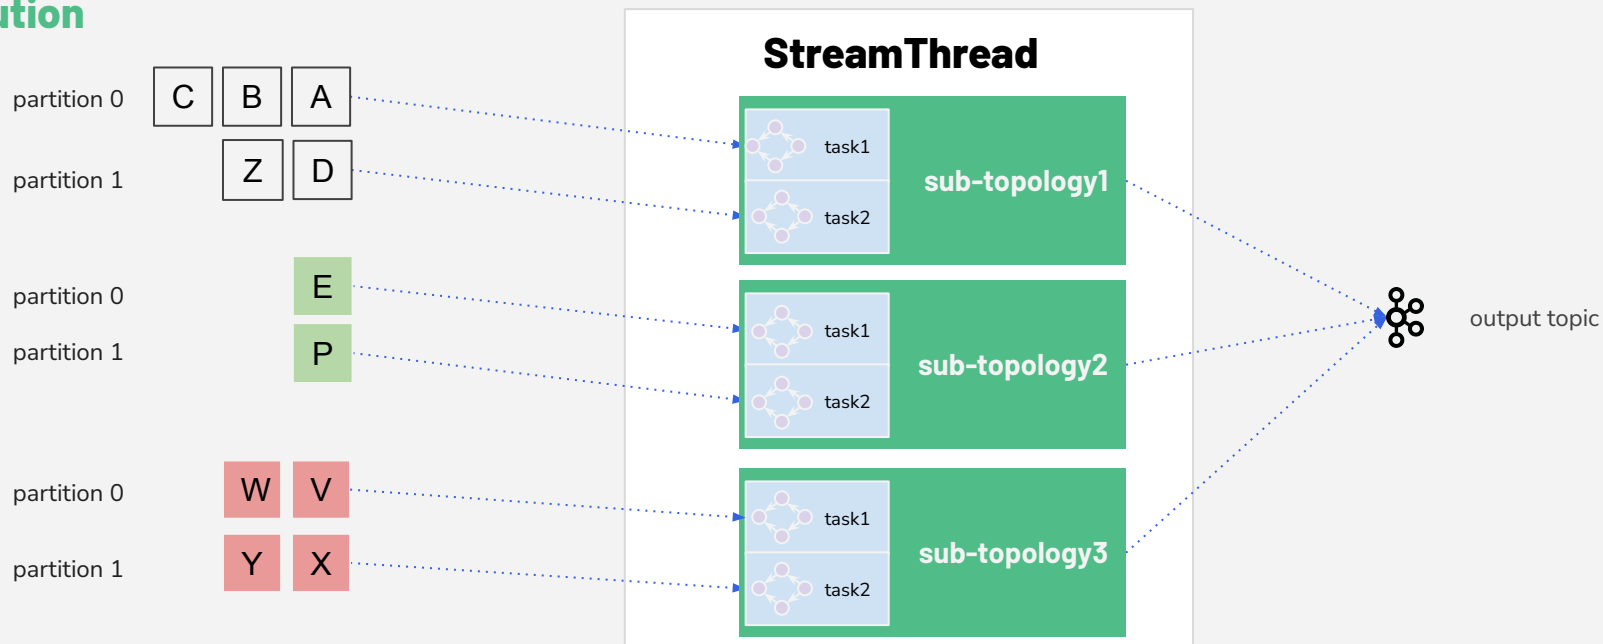
## Additional partitions and tasks may improve throughput






"A sub-topology is a set of processors, that are all transitively connected as parent/child or via state stores in the topology, so different sub-topologies exchange data via topics and don't share any state stores. Each task may instantiate only one such sub-topology for processing. This further scales out the computational workload to multiple tasks." <sup>1</sup>

## Solution

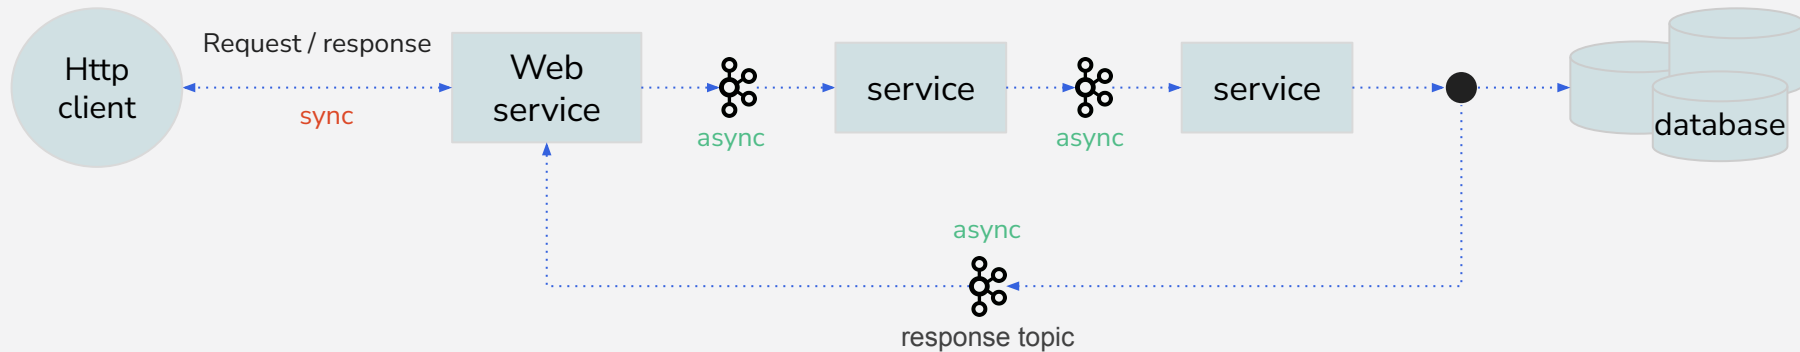


# 5. Broadcasting Events



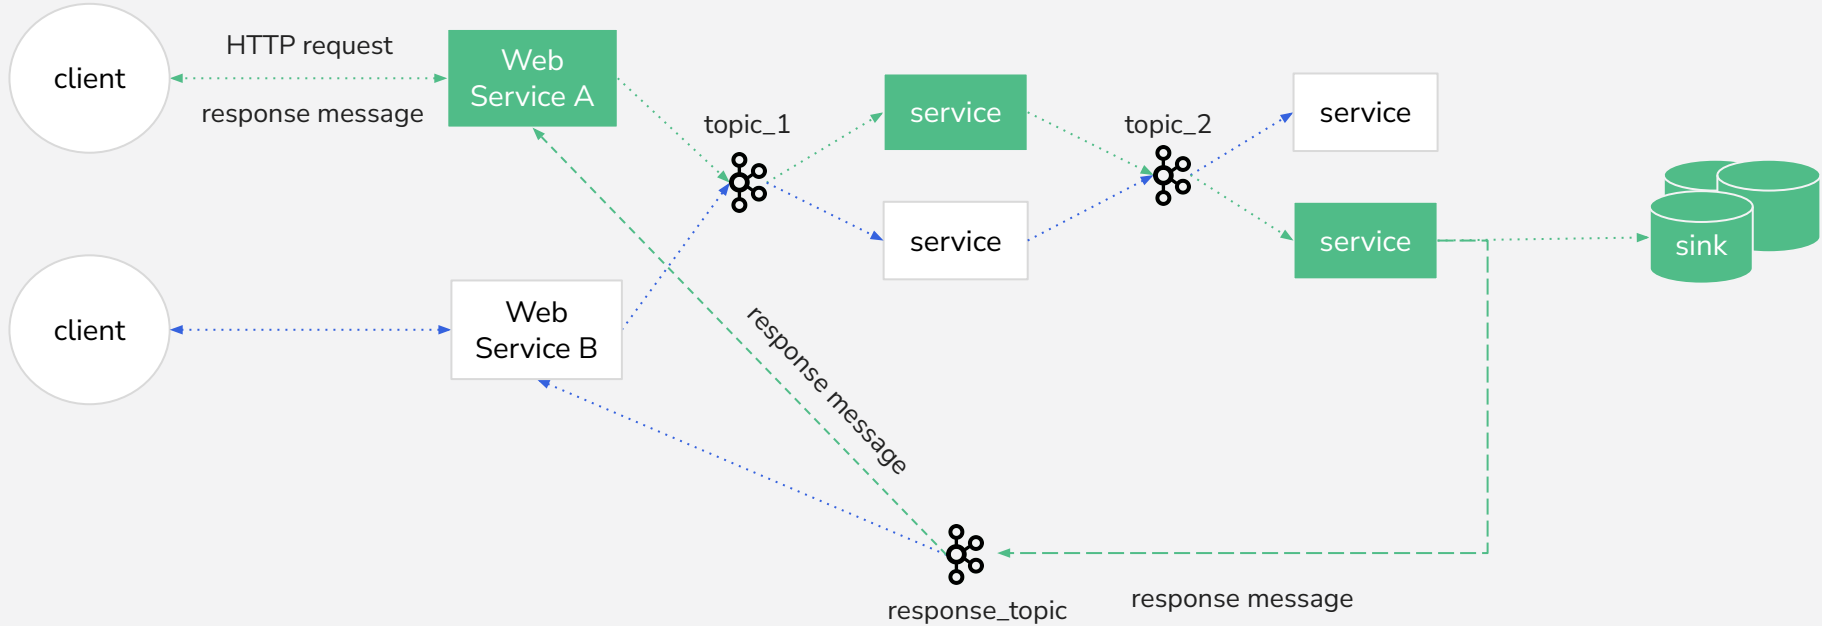
## Problem 1

Kafka asynchronous message event stream is not suited to synchronous message handling (e.g. HTTP)



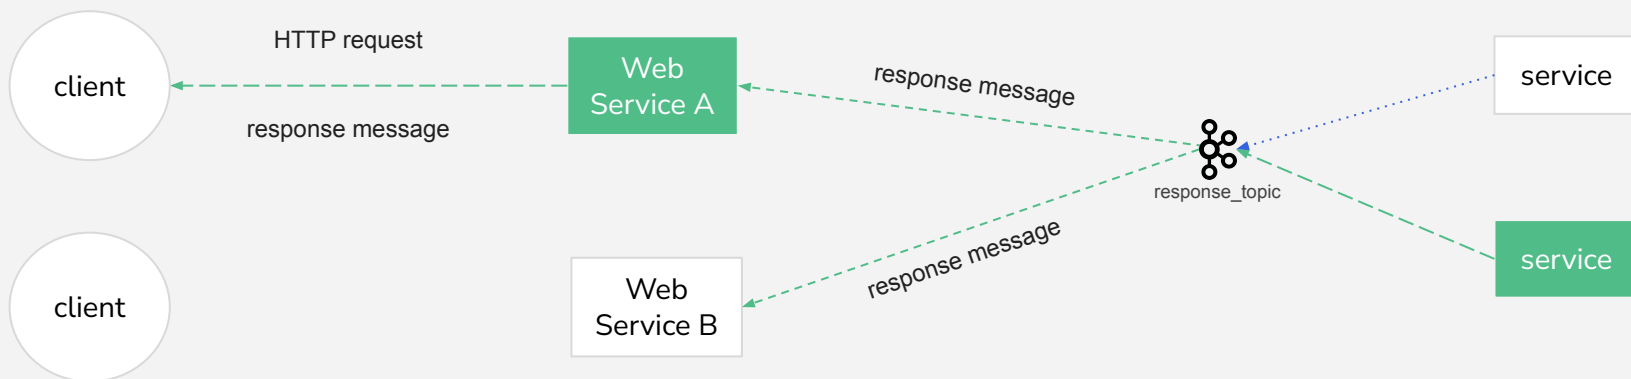
## Problem 2

Request and response traverse an unpredictable path across service nodes



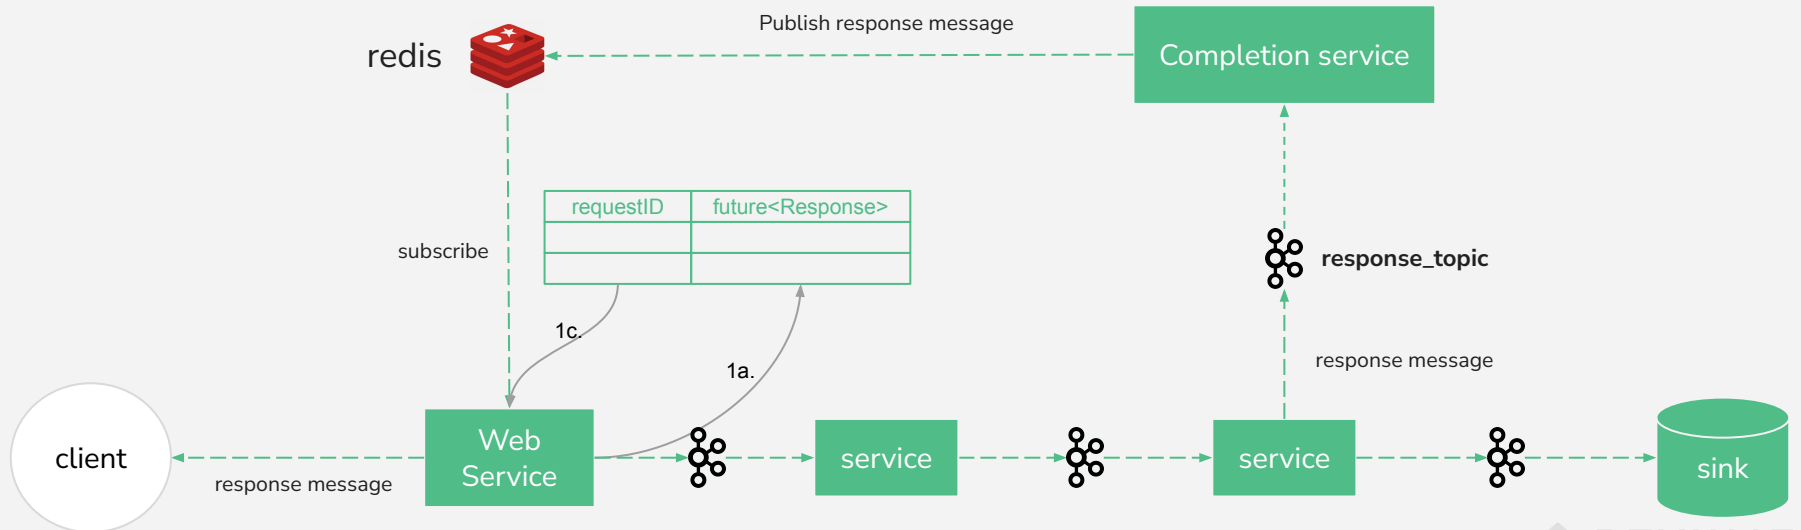
## Goal

Ensure that the same (webservice) that received the client HTTP request message will receive the response generated in the asynchronous service pipeline




## Solution

1. Webservice listener thread
  - a. Enter `<requestId, CompletableFuture<Response>>` in `responseTable`
  - b. Post request to asynchronous pipeline
  - c. Block on future
2. Webservice redis channel subscriber thread
  - a. Listen for all responses
  - b. Look up `requestId` in `responseTable`
  - c. Complete future if found



## 6. Processing Messages in Parallel (Beyond Partitions)



*"...the maximum parallelism at which your application may run is bounded by the maximum number of stream tasks, which itself is determined by maximum number of partitions of the input topic(s) the application is reading from" <sup>1</sup>*

## Latency vs. Throughput

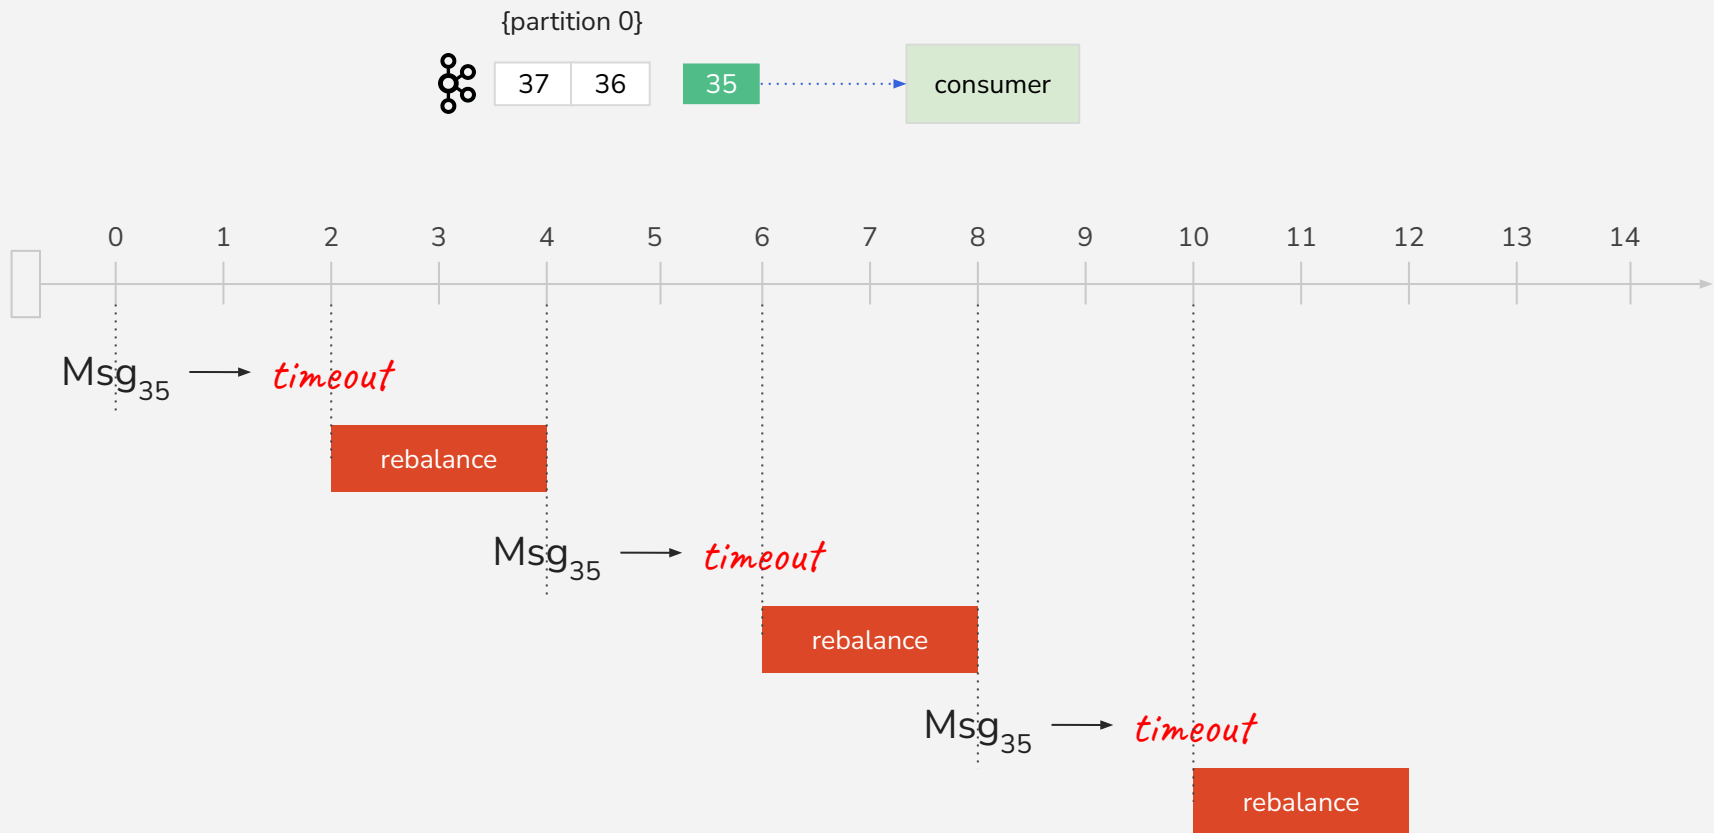
- Latency = seconds per message = delay
- Throughput = messages per second (service overall) = volume

## Problem

- Fixed number of topic partitions
- Message latency is high
  - Consumer lag
  - Polling timeouts resulting in rebalance events
    - “Stuck partition” (a.k.a. “slow consumer”)



# WHAT IS A "STUCK PARTITION"?



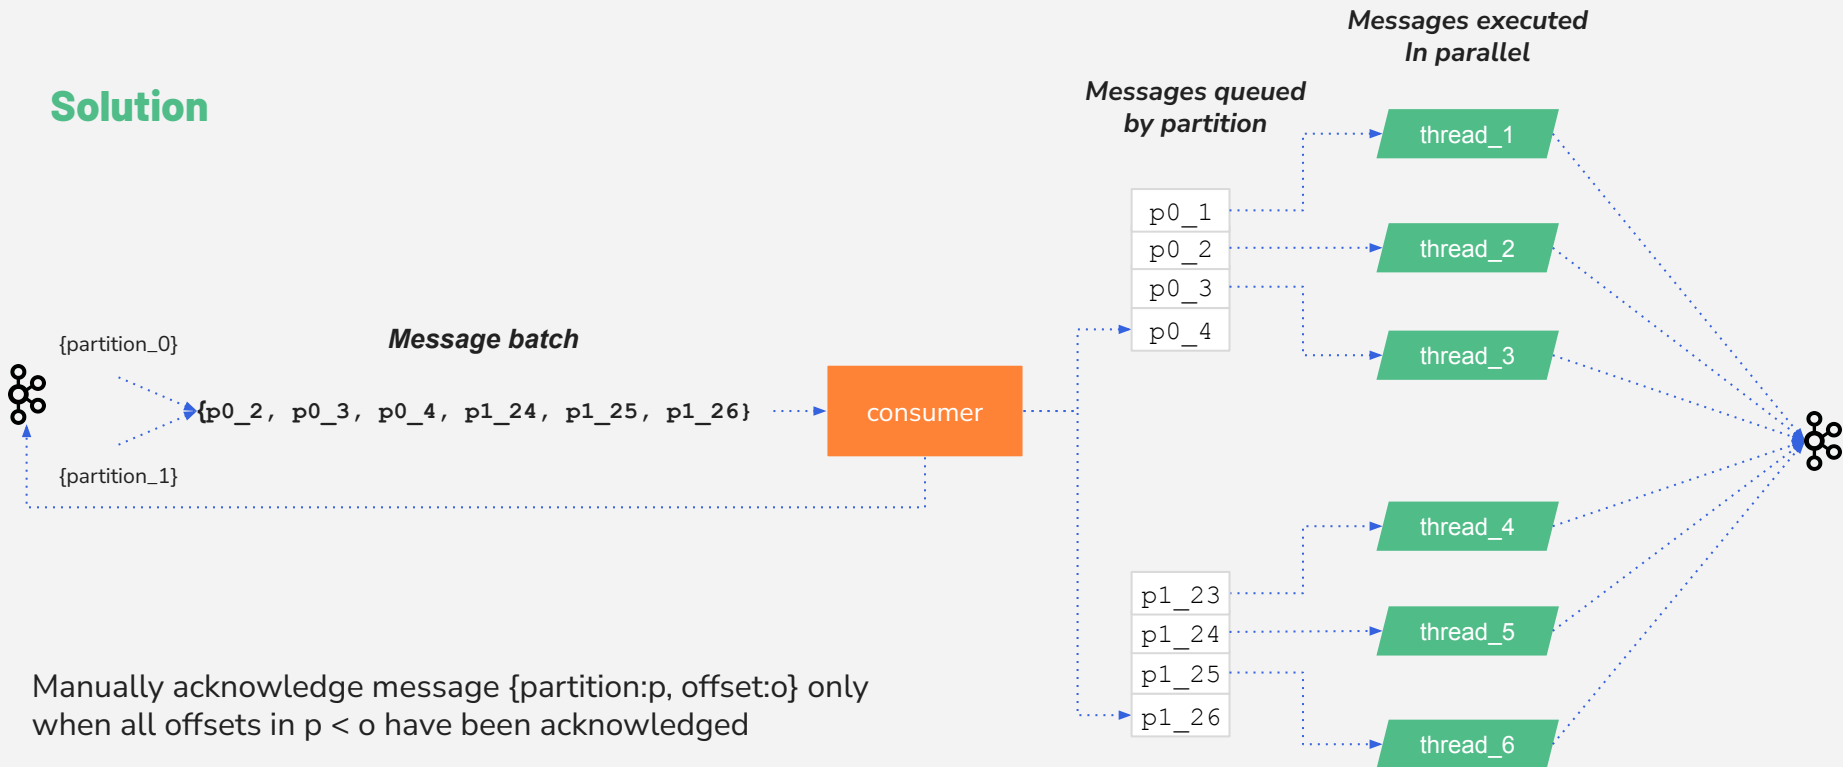
## Goal

- Increase message throughput (msgs/sec)
  - Without increasing number of partitions
  - Without improving message latency
- Prevent stuck partitions
- All messages are processed
- Message scheduling order is preserved

## Caveat

- Message processing order is not guaranteed

## Solution



Manually acknowledge message {partition:p, offset:o} only when all offsets in  $p < o$  have been acknowledged

## Some things to consider

- Kafka vs. Kafka Streams?
- Is process latency consistent for all messages?
- Are their (rare) outlier messages with much higher latency than average?
- What is cause of high message latency?
  - Processing complexity? (CPU bound)
  - Synchronous dependencies on other services? (blocking)
- Use the Confluent Parallel Consumer? <sup>1</sup>
- Wait for KIP-932 (Queues for Kafka)? <sup>2</sup>



**Questions?**